

VISHAL SHARMA

Bangalore, India | +91-6396628186 | iamvishal.inbox@gmail.com | [LinkedIn](#) | [Portfolio](#)

SUMMARY

R&D Lead owning all AI product work at a high-throughput operation. Shipped **5 production systems** solo: a real-time **voice agent** under 300ms latency, a **7-agent autonomous OS** at \$55-95/month infra, **RAG support agents** at **89.6% self-service resolution**, a **public MCP server**, and a revenue platform with **\$15K+ in 7 days**. Strong at agentic orchestration, cost-per-outcome thinking, and taking systems from zero to daily production use.

SKILLS

Languages: Python, TypeScript, JavaScript, Go, SQL, Java, Embedded C

AI / LLM: LLM APIs (Anthropic, OpenAI, Gemini), RAG, FAISS, MCP (Model Context Protocol), LangChain-style orchestration, prompt engineering, Pydantic structured output, multi-agent systems

Frameworks: FastAPI, asyncio, Next.js, React, Node.js, discord.py, REST API, Tailwind CSS

Real-time and Voice: LiveKit (WebRTC), ElevenLabs TTS/STT

Databases: PostgreSQL (Supabase, RLS), Redis, SQLite, Cloudflare KV, FAISS (vector)

Infrastructure: Docker, Cloudflare Workers, Vercel Edge, Tailscale, N8N, Git and CI/CD

Tools and APIs: PostHog, Razorpay, Google Drive API, Discord API, Resend, Postman

EXPERIENCE

R&D Lead (AI Engineer) | 100xEngineers

Aug 2025 – Present | Bangalore, India

Sole engineer on all AI product work. Report to CTO. Five production systems built end-to-end, detailed below.

Zeno - Voice AI teaching system

- Built the full voice pipeline: WebRTC transport, streaming STT/TTS, RAG retrieval over 80+ docs. Held end-to-end **latency under 300ms** across concurrent live sessions. Runs in daily production for mentors.
- Designed a **Socratic** system prompt that guides mentors through 6 to 7 structured turns before generating curriculum. Red-teamed against hallucinations and turn-taking failures across 30+ live sessions.
- Cut** per-session **cost 80% (\$1+ to under \$0.20)** by separating LLM reasoning from deterministic rendering: model produces a typed JSON spec, Python handles output with zero tokens in the render step. Architectural decision, not a prompt trick.

EDITH - Multi-agent company OS

- Engineered **7+ specialized agents** with hard trust boundaries per role: draft-only vs. autonomous execution, isolated identity and memory files per agent. Caught context-bleed and tone failures across **20+ test runs** before any output reached live workflows.
- Reclaimed **30+ hours/month** of leadership bandwidth at **\$55 to \$95** monthly all-in infra. Pipeline runs research, content drafting, email triage, and candidate screening without human triggers.

Sage and Scout - Conversational support bots

- Two RAG agents handling **1,000+ queries/week** at **89.6% self-service resolution**. Reduced human handling time by **30+ hours/month**. Unresolved queries auto-routed via 2-layer logic: keyword match at near-zero latency, RAG score fallback for ambiguous cases.

Wiki Plugin Store - Open-source AI infrastructure

- Built and **open-sourced** a **5+ production MCP server** exposing an enterprise knowledge base inside any MCP-compatible client. Extensible plugin architecture with modular skills and connectors.

Community Events Platform - Ticketing and lead capture

- Solo-built a high-load transactional platform. Generated **\$15,000+ in 7 days**, served **100,000+ visitors** at zero downtime. Registration flow designed as the primary lead acquisition funnel.

Hardware Systems Intern | Schneider Electric

Mysuru, India

- Improved production quality **10%** with a 5-sensor ultrasonic welding error detection system. Built a GPS-enabled smart lock for shipment tracking (patent filing in progress).

EDUCATION

B.Tech. Electronics and Communication, VIT Chennai | GPA 8.58 / 10